

Project Overview

What Code Archaeologist is, what it does, and where everything else lives.

Read this if: everyone on the team, first thing

Code Archaeologist is an AI agent that explains not just what code does, but why it exists. It reads a public GitHub repository's code, history, and structure, and answers questions with cited evidence instead of guesses.

What it does

- Analyzes a public GitHub repository (structure, history, contributors, code entities)
- Answers "why" questions about the code with evidence and a confidence level
- Gives task-specific guidance on where to make a change and why
- Flags fragile or high-risk areas using deterministic signals, not guesswork
- Captures human explanations when the evidence alone isn't enough, and reuses them later
- Runs an offboarding analysis to capture a leaving contributor's knowledge before it's lost

What it will not do

- Execute any code from the analyzed repository
- Work with private repositories
- Answer questions unrelated to the repository being analyzed

Document set

Document	What's in it
PRD.md	Product requirements — what we're building and for whom
MASTER_BUILD_SPEC.md	Technical scope, stack, and build priorities
API_CONTRACT.md	Frontend/backend interface — endpoints and payloads
DATABASE_DESIGN.md	Shared database schema
TEAM_WORKFLOW.md	Git process and integration order
TEAM_ASSIGNMENTS.md	Who owns what
CONTRIBUTING.md	Rules for opening a PR

Tech stack

Backend: Python, FastAPI, PostgreSQL, pgvector, LangGraph, Groq.
Frontend: React, TypeScript, Vite, Tailwind CSS, Monaco Editor, React Flow.
Infra: Supabase (auth + database), Render (backend), Vercel (frontend).